



University of
New Haven

TAGLIATELA
COLLEGE OF ENGINEERING

MASTER OF SCIENCE IN DATA SCIENCE

Deep Learning Final
Project Presentation
[Working Link](#)





Leaf Health Semantic Segmentation Using SegFormer

Meet the Team



Lakshman Rajith




Vittu Darshan



Yogananda Manjunath

Introduction

- Task Description: The goal is to perform semantic segmentation of leaf regions into background, healthy leaf and dry leaf
- Target Object: Pixel-level classification of leaf health to support plant disease detection and agricultural automation.
- Why This Task Matters?
 - Manual inspection is slow and unreliable.
 - Automated segmentation helps monitor plant stress early.
 - Applicable in smart farming and environmental monitoring.





Dataset Overview

Collected 200 real leaf images using mobile phone camera

Each image has a manually created pixel level mask

Classes:

1. 0 – background
2. 1 – Healthy Leaf
3. 2 – Dry Leaf

Masks created using CVAT

Resolution standardized to 512x512



Healthy Leaf



Dry Leaf




Dataset Details

- Task
 - Multi-class semantic segmentation
 - Predict health category for each pixel
- Dataset Link: [Dataset Link](#)
- Number of Samples
 - Total Images: ~200
 - Train: 80%
 - Validation: 10%
 - Test: 10%
- Image Size: 512×512 (after preprocessing)
- Normalization
 - Mean = (0.485, 0.456, 0.406)
 - Std = (0.229, 0.224, 0.225)
- Data Augmentation: Affine transforms, random brightness/contrast, grid distortion, elastic transform, noise + motion blur, horizontal & vertical flips and coarse dropout

Dataset Annotations

Annotation Structure (1 example)

```
{  
    "image": "161.png",  
    "mask": "161.png",  
    "classes": {  
        "0": "background",  
        "1": "healthy_leaf",  
        "2": "dry_leaf"  
    }  
}
```

Leaf No	Normal Image	Masked Image
10		
35		




Methods (Overview)

- Model Used: SegFormer-B0
- Approach:
 - Use pretrained SegFormer encoder (MiT) from ADE20K
 - Replace segmentation head with 3-class output
 - Finetune using your leaf dataset
 - Improve performance with augmentation + Dice Loss
- Training Strategy:
 - CE finetuning → evaluate
 - CE + Dice finetuning → compare
 - Report both qualitative & quantitative results

Model Architecture

Base Neural Network: SegFormer-B0	Modifications to Adapt to this Task
<ul style="list-style-type: none">• Overlapping Patch Embedding• 4-stage MiT Transformer Encoder• Multi-scale features ($H/4 \rightarrow H/32$)• Lightweight MLP Decoder for segmentation• Final upsampling to full resolution	<ul style="list-style-type: none">• Changed head to output 3 classes• Applied strong augmentations• Introduced Dice Loss to handle class imbalance• Tuned optimizer and LR via random search





Loss Functions

Training Objective Functions

- Cross Entropy Loss
 - Standard multi-class segmentation loss
- Dice Loss (Recent Technique)
 - Improves class imbalance handling
 - Enhances boundary segmentation
- Combined Loss = $0.5 \text{ CE} + 0.5 \text{ Dice}$



Experiments

Experiment Pipeline

Baseline inference using pretrained weights

CE finetuning for 10 epochs

CE evaluation (IoU, confusion matrix, samples)

CE + Dice finetuning

Comparison of improvements

Metrics Used


Pixel Accuracy

Mean IoU

Confusion Matrix

Qualitative visualizations

Qualitative Results Before Training





Optimization & Hyperparameter Tuning

Optimizers Tried: AdamW, SGD

Best Optimizer: AdamW

Hyperparameters Tuned

- Learning Rate: $[5e^{-5} \rightarrow 5e^{-4}]$
- Batch Size: [4, 8, 16]
- Weight Decay: [0.0, 0.05]

Tuning Strategy

- Random Search over 20 combinations
- Evaluate based on validation IoU

Best Parameters Found

- LR = 1e-4
- Batch = 8
- Weight Decay = 0.05
- Optimizer = AdamW



Finetuning Results (CE)

Experiment Setup



- CE-only finetuning
- 10 epochs
- Batch size 8
- AdamW optimizer

Quantitative Results

- Val IoU ≈ 0.83
- Good segmentation of healthy & dry leaf
- Strong generalization



Results After Training (CE)





Transfer Learning vs Pretrained Model

Pretrained (Before Training)

- Incorrect segmentation
- High confusion with background
- Fails to separate healthy/dry leaf regions

After Fine-tuning

- Clear boundaries
- Accurate leaf region segmentation
- Strong detection of dry leaf patches

Conclusion:

- Transfer learning dramatically improves performance.



Dice Loss

What is Dice Loss?

- Measures the overlap between predicted and true regions.
- Designed for segmentation tasks with class imbalance.

Why Use It?

- Improves detection of small or rare regions (like dry leaf tissue).
- Produces sharper boundaries and better region consistency.


Formula:

$$\text{Dice} = 2\text{TP} / (2\text{TP} + \text{FP} + \text{FN})$$


Results After Dice Loss (Improved Technique)

CE + Dice Fine-tuning

- Better dry leaf localization
- Fewer false negatives
- Improved boundary accuracy
- Higher IoU for minority class



Qualitative Results





Conclusion

Key Takeaways

- SegFormer effectively segments leaf health regions
- Strong results with small dataset
- Hyperparameter tuning improved baseline
- Dice Loss significantly improved dry-leaf detection
- Model can be extended to other plant diseases



References

- [1] E. Xie, W. Wang, Z. Yu, A. Anandkumar, J. M. Alvarez, and P. Luo, "SegFormer: Simple and Efficient Design for Semantic Segmentation with Transformers," *arXiv preprint arXiv:2105.15203*, 2021. <https://arxiv.org/abs/2105.15203>
- [2] F. Milletari, N. Navab, and S. A. Ahmadi, "V-Net: Fully Convolutional Neural Networks for Volumetric Medical Image Segmentation," *arXiv preprint arXiv:1606.04797*, 2016. <https://arxiv.org/abs/1606.04797>
- [3] J. Long, E. Shelhamer, and T. Darrell, "Fully Convolutional Networks for Semantic Segmentation," *arXiv preprint arXiv:1411.4038*, 2015. <https://arxiv.org/abs/1411.4038>
- [4] A. Dosovitskiy *et al.*, "An Image is Worth 16×16 Words: Transformers for Image Recognition at Scale," *arXiv preprint arXiv:2010.11929*, 2020. <https://arxiv.org/abs/2010.11929>
- [5] A. Buslaev, V. I. Iglovikov, E. Khvedchenya, A. Parinov, M. Druzhinin, and A. A. Kalinin, "Albumentations: Fast and Flexible Image Augmentations," *Information*, vol. 11, no. 2, 2020. <https://www.mdpi.com/2078-2489/11/2/125>
- [6] I. Loshchilov and F. Hutter, "Decoupled Weight Decay Regularization," *arXiv preprint arXiv:1711.05101*, 2019. <https://arxiv.org/abs/1711.05101>



Thank you!